

Assignment -4

Python Programming

Assignment Date	30 October 2022
Student Name	R.Vigneshwari
Student Roll Number	820419106070
Maximum Marks	2 Marks

1. Download the dataset

2. Load the dataset into the tool.

```
In [1]: import pandas as pd
data = pd.read_csv("Mall_Customers.csv")
data.head()
```

```
Out[1]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
In [2]: data.shape
```


```
Out[2]: (200, 5)
```

3. Perform Below Visualizations

· Univariate Analysis

```
In [3]: import seaborn as sns
sns.histplot(data, x="Age")
```


```
Out[3]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```



In []: · Bi- Variate Analysis

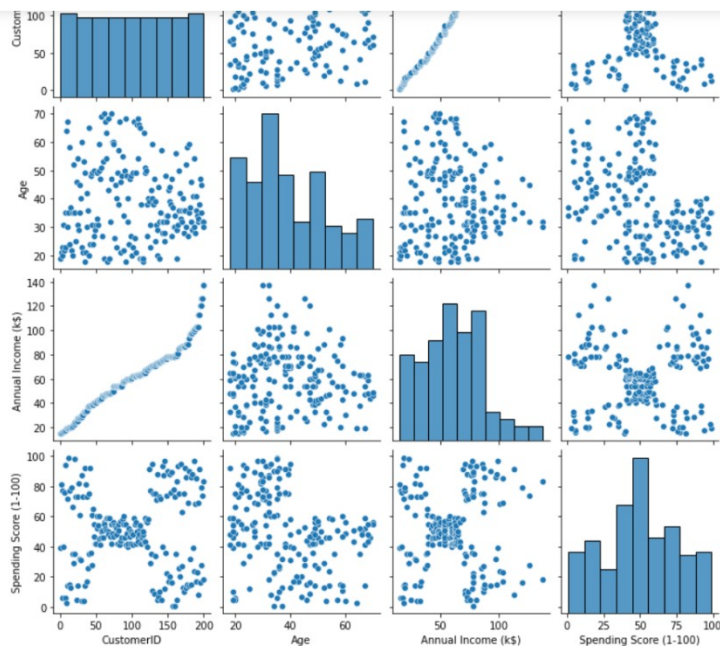
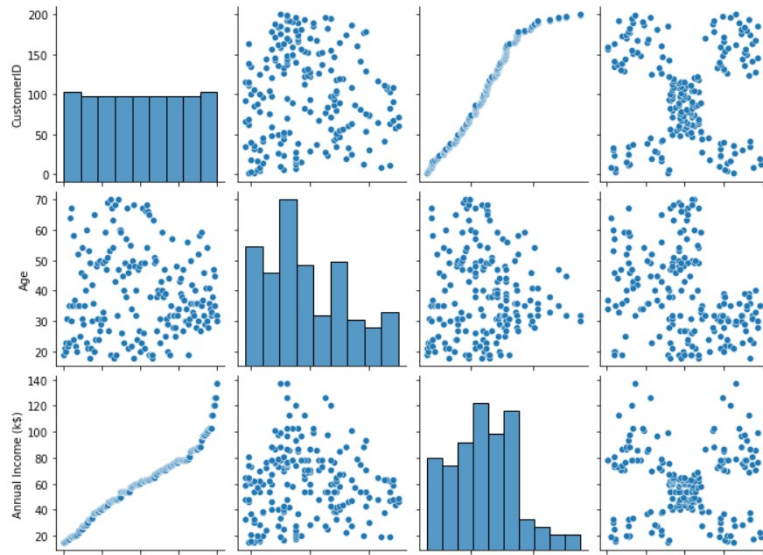
```
In [4]: sns.displot(data, x="Age", y="Annual Income (k$)")
```

```
Out[4]: <seaborn.axisgrid.FacetGrid at 0x1c948fa0370>
```



In [6]: `sns.pairplot(data)`

Out[6]: `<seaborn.axisgrid.PairGrid at 0x1c94373c6d0>`



4. Perform descriptive statistics on the dataset

In [7]: `data.mean()`

C:\Users\welcome\AppData\Local\Temp\ipykernel_13532\531903386.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
`data.mean()`

Out[7]: CustomerID 100.50
Age 38.85
Annual Income (k\$) 60.56
Spending Score (1-100) 50.20
dtype: float64

In [8]: `data.median()`

C:\Users\welcome\AppData\Local\Temp\ipykernel_13532\4184645713.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
`data.median()`

Out[8]: CustomerID 100.5
Age 36.0
Annual Income (k\$) 61.5
Spending Score (1-100) 50.0
dtype: float64

```
In [9]: data.mode()
```

```
Out[9]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Female	32.0	54.0	42.0
1	2	NaN	NaN	78.0	NaN
2	3	NaN	NaN	NaN	NaN
3	4	NaN	NaN	NaN	NaN
4	5	NaN	NaN	NaN	NaN
...
195	196	NaN	NaN	NaN	NaN
196	197	NaN	NaN	NaN	NaN
197	198	NaN	NaN	NaN	NaN
198	199	NaN	NaN	NaN	NaN
199	200	NaN	NaN	NaN	NaN

200 rows × 5 columns

```
In [10]: data.skew()
```

C:\Users\welcome\AppData\Local\Temp\ipykernel_13532\1188251951.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
data.skew()

```
Out[10]: CustomerID      0.000000  
Age      0.485569  
Annual Income (k$)      0.321843  
Spending Score (1-100) -0.047220  
dtype: float64
```

```
In [11]: data.kurt()
```

C:\Users\welcome\AppData\Local\Temp\ipykernel_13532\2907027414.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
data.kurt()

```
Out[11]: CustomerID      -1.200000  
Age      -0.671573  
Annual Income (k$)      -0.098487  
Spending Score (1-100) -0.826629  
dtype: float64
```

```
In [12]: data.std()
```

C:\Users\welcome\AppData\Local\Temp\ipykernel_13532\2723740006.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
data.std()

```
Out[12]: CustomerID      57.879185  
Age      13.969007  
Annual Income (k$)      26.264721  
Spending Score (1-100)  25.823522  
dtype: float64
```

```
In [13]: data.var()
```

C:\Users\welcome\AppData\Local\Temp\ipykernel_13532\445316826.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
data.var()

```
Out[13]: CustomerID      3350.000000  
Age      195.133166  
Annual Income (k$)      689.835578  
Spending Score (1-100)  666.854271  
dtype: float64
```

5. Check for Missing values and deal with them

```
In [14]: data.isna().sum()
```


```
Out[14]: CustomerID      0  
Gender      0  
Age      0  
Annual Income (k$)      0  
Spending Score (1-100)  0  
dtype: int64
```


6. Find the outliers and replace them outliers

```
In [15]: sns.boxplot(data['Annual Income (k$)'])
```

C:\Users\welcome\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

```
Out[15]: <AxesSubplot:xlabel='Annual Income (k$)'\>
```






```
In [16]: import numpy as np
data['Annual Income (k$)']=np.where(data['Annual Income (k$)']>120,20,data['Annual Income (k$)']) #replacing
```

```
In [17]: sns.boxplot(data['Annual Income (k$)'])
```

C:\Users\welcome\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
Out[17]: <AxesSubplot:xlabel='Annual Income (k$)'\>
```



7. Check for Categorical columns and perform encoding

```
In [18]: data.head()
```

```
Out[18]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
In [19]: from sklearn.preprocessing import LabelEncoder
```

```
In [20]: le=LabelEncoder()
```

```
In [21]: data['Gender']=le.fit_transform(data['Gender'])
data.head()
```

```
Out[21]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6
3	4	0	23	16	77
4	5	0	31	17	40

8. Scaling the data

```
In [22]: from sklearn.preprocessing import StandardScaler
```

```
In [23]: scaler=StandardScaler()
x=scaler.fit_transform(data)
```

```
In [24]: print(data)
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	1	19	15	39
1	2	1	21	15	81
2	3	0	20	16	6
3	4	0	23	16	77
4	5	0	31	17	40
..
195	196	0	35	120	79
196	197	0	45	20	28
197	198	1	32	20	74
198	199	1	32	20	18
199	200	1	30	20	83

```
[200 rows x 5 columns]
```

9. Perform any of the clustering algorithms


```
In [25]: import matplotlib.pyplot as plt
         from sklearn.cluster import KMeans

In [26]: x = data.iloc[:, [3, 4]].values

In [27]: print(x.shape)
(200, 2)

In [28]: km = KMeans(n_clusters = 5, init = 'k-means++', max_iter = 300, n_init = 10, random_state = 0)
         y_means = km.fit_predict(x)

In [29]: plt.scatter(x[y_means == 0, 0], x[y_means == 0, 1], s = 100, c = 'pink', label = 'miser')
         plt.scatter(x[y_means == 1, 0], x[y_means == 1, 1], s = 100, c = 'yellow', label = 'general')
         plt.scatter(x[y_means == 2, 0], x[y_means == 2, 1], s = 100, c = 'blue', label = 'target')
         plt.scatter(x[y_means == 3, 0], x[y_means == 3, 1], s = 100, c = 'magenta', label = 'spendthrift')
         plt.scatter(x[y_means == 4, 0], x[y_means == 4, 1], s = 100, c = 'orange', label = 'careful')
         plt.scatter(km.cluster_centers_[0,0], km.cluster_centers_[0, 1], s = 50, c = 'black', label = 'centroid')
         plt.style.use('fivethirtyeight')
         plt.title('K Means Clustering', fontsize = 20)
         plt.xlabel('Annual Income')
         plt.ylabel('Spending Score')
         plt.legend()
         plt.grid()
         plt.show()
```



10. Add the cluster data with the primary dataset

```
In [30]: from sklearn.cluster import KMeans

In [31]: km=KMeans(n_clusters=3, random_state=0)

In [32]: data['Group or Cluster'] = km.fit_predict(data)

In [34]: data.head()

Out[34]:
```

CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Group or Cluster	
0	1	1	19	15	39	2
1	2	1	21	15	81	2
2	3	0	20	16	6	2
3	4	0	23	16	77	2
4	5	0	31	17	40	2

11. Split the data into dependent and independent variables.

```
In [35]: y=data['Spending Score (1-100)']
         y.head()

Out[35]:
```

0	39
1	81
2	6
3	77
4	40

Name: Spending Score (1-100), dtype: int64

```
In [36]: x=data.drop(columns=['Spending Score (1-100)'])
         x.head()

Out[36]:
```

CustomerID	Gender	Age	Annual Income (k\$)	Group or Cluster	
0	1	1	19	15	2
1	2	1	21	15	2
2	3	0	20	16	2
3	4	0	23	16	2
4	5	0	31	17	2

12. Split the data into training and testing

```
In [37]: from sklearn.model_selection import train_test_split

In [38]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)

In [39]: print(x_train.shape,y_train.shape)

(160, 5) (160,)

In [40]: print(x_test.shape,y_test.shape)

(40, 5) (40,)
```

13. Build the Model

```
In [41]: from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier()
```

14. Train the Model

```
In [42]: model.fit(x_train,y_train)

Out[42]: DecisionTreeClassifier()
```

15. Test the Model

```
In [43]: pred2=model.predict(x_test)

pred2

Out[43]: array([[49, 13,  4, 90,  5, 54, 42, 95, 35, 81, 42, 16, 27, 55, 35, 89, 45,
16, 52, 55, 40, 99, 72, 42, 55, 46, 55, 55, 71, 40, 99, 15, 60, 55,
16, 55, 40, 71, 16, 42], dtype=int64)
```

16. Measure the performance using Evaluation Metrics.

```
In [44]: from sklearn import metrics

In [45]: metrics.confusion_matrix(y_test,pred2)

Out[45]: array([[0, 0, 0, ..., 1, 0, 0],
[0, 1, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]], dtype=int64)

In [46]: print('DT model ACcuracy Score:',metrics.accuracy_score(y_test,pred2))

DT model ACcuracy Score: 0.025
```

```
In [45]: metrics.confusion_matrix(y_test,pred2)

Out[45]: array([[0, 0, 0, ..., 1, 0, 0],
[0, 1, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]], dtype=int64)

In [46]: print('DT model ACcuracy Score:',metrics.accuracy_score(y_test,pred2))

DT model ACcuracy Score: 0.025
```

```
In [47]: acc=metrics.accuracy_score(y_test,pred2)
acc

Out[47]: 0.025
```

```
In [48]: 1-acc

Out[48]: 0.975
```

```
In [49]: data.head()

Out[49]:
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	Group or Cluster
0	1	1	19	15	39	2
1	2	1	21	15	81	2
2	3	0	20	16	6	2
3	4	0	23	16	77	2
4	5	0	31	17	40	2